

Machine Learning and Data Mining project: Leaf Classification

Michele Alessi¹, Samuele D'Avenia², and Elena Rivaroli³

^{1, 2, 3} problem statement, solution design, solution development,
data gathering, writing

Course of AA 2022-2023 - Data Science and Scientific Computing

1 Problem statement

Leaf classification is the problem of assigning the correct tree species to an unknown observed leaf. Traditionally, this was the job of experts called taxonomists. Through the usage of modern automatic recognition systems, it is possible to develop an efficient tool to help with this classification process. If some examples, i.e. leaves with the correct species assigned, are provided, this can easily be seen as a supervised machine learning problem.

The purpose of this project is to propose a leaf classifier after comparing different techniques. This is done using a dataset of leaves containing some of their attributes, which are further described below.

2 Assessment and performance indexes

The effectiveness index was selected by taking into account that this is multi-class problem and the dataset is not perfectly balanced. The chosen assessment measure is *balanced accuracy* bA , which is defined as follows [2]:

$$bA = \frac{\sum_{i=1}^n \frac{|TC_i|}{|C_i|}}{n}$$

where C_i , $i = 1, \dots, n$ are the n classes in the dataset D and TC_i is the set of observations correctly predicted into C_i , $i = 1, \dots, n$. Since the True Positive Rate (TPR) for each class is computed disregarding other classes, it's clear that bA gives the same weight to each class; this means that smaller classes have the same influence on the final effectiveness as the biggest one; in particular, bA is insensitive to imbalanced class distribution. It is important to report the existence of other effectiveness indexes suitable for multi-class classification problems with imbalanced classes, which have been analyzed as well [3]. The

final choice fell onto the balanced accuracy for explainability, interpretability and efficiency reasons. It is fast to compute and on equal terms, it is the simplest index to interpret.

3 Proposed solution

The proposed solution uses a nested cross-validation approach, which is recommended when dealing with small datasets [4]. The outer cross-validation is used to compare the different models, using an appropriate effectiveness measure to choose the best one (i.e. performance index described in Section 2). The internal nested cross-validation is used to tune the hyper-parameters using the same effectiveness measure (more details on the choice of hyper-parameters are provided in section 4.2) through a grid-search approach.

Since the dataset is not large and it is slightly imbalanced, both cross-validations are performed using stratification, to obtain better estimates [4]. The number of folds in the inner and outer loop is chosen to ensure the presence of at least one observation of each class in each inner fold.

The description provided below reports the hyper-parameters tuning procedure of a learning technique $(f'_{\text{learn}}, f'_{\text{predict}})$ and to obtain the effectiveness measure v_{effect} to compare the various effectiveness techniques and select the most appropriate one.

1. Split the whole dataset *data* into k_{outer} folds using stratified splitting.
2. For $i = 1, 2, \dots, k_{\text{outer}}$:
 - (a) Consider all folds except fold i . Denote the subset of the data containing all folds but i as $data_{(-i)}$.
 - (b) Select a finite set of values for each hyper-parameter.
 - (c) In order to perform grid search, for each possible combination of hyper-parameters p repeat:
 - i. For $j = 1, 2, \dots, k_{\text{inner}}$:
 - A. Consider all folds except fold j . Denote the subset of $data_{(-i)}$ containing all folds but j as $data_{(-i, -j)}$.
 - B. Learn a model on $data_{(-i, -j)}$ using the hyper-parameters p of the current iteration.
 - C. Obtain the effectiveness index on fold j , denote it as v_j .
 - ii. With this set of hyper-parameters p compute the average effectiveness denoted as $v_p = \text{mean}_j(v_j)$.
 - (d) Choose as best hyper-parameters the combination which produced the largest v_p .
 - (e) Train the model with the best hyper-parameters on $data_{(-i)}$.
 - (f) Compute $v_{\text{eff}, i}$, the effectiveness measure obtained by applying the learnt prediction function on fold i .

3. Compute v_{eff} as mean of the effectiveness measures obtained $v_{\text{eff}, i}$.

The learning technique with the highest v_{effect} is chosen as the best one. The final step is to choose the best hyper-parameters for the selected technique using a simple cross-validation on the whole dataset. The v_{effect} obtained is used as effectiveness measure of the proposed classifier.

4 Experimental evaluation

4.1 Data

The leaf dataset used for this task comprises 340 leaf observations from 30 different plant species. For each leaf, *Class*, *Specimen Number* and 14 different other features are available. Variable *Class* contains the plant species of each leaf and is used as the label of the supervised learning technique. Further description of the variables can be found at [5].

4.2 Procedure

The variable *Specimen Number* is removed from the dataset as it is only used as an identifier for each leaf within each class. The dataset was prepared with the task of leaf classification in mind, as such no further feature engineering is deemed necessary.

The techniques tried are decision tree, random forest (RF), k-nearest neighbours (kNN) and support vector machines (SVM) with Gaussian and linear kernel as they are some of the most widely used. For the last three techniques (the two SVM and kNN), standardization is performed to ensure all dependent variables are on the same scale. It is performed in each inner fold of the procedure to avoid the algorithm having knowledge of unseen data (i.e. left out fold data). Moreover, in order to allow SVM to work with the multi-class problem at hand, a one-vs-all approach is taken for efficiency reasons.

For RF, the default number of variables to consider at each split (\sqrt{p} where p is the number of variables) is deemed satisfactory. As such no tuning is performed for that hyper-parameter. For similar reasons, the Gini criterion is used as node impurity measure for both the decision tree and the RF [6].

Table 1 below reports what hyper-parameters are included in the grid search procedure (along with a brief description) for each technique. It also contains the results of the nested cross-validation procedure (outlined in Section 3) in the last two columns. These two represent the mean effectiveness score v_{effect} and the standard deviation (s.d.). It appears that the RF and SVM with Gaussian and linear kernel algorithms perform slightly better. In fact, they have a larger mean balanced accuracy. However, RF seems to have a smaller standard deviation. It also requires tuning of only one hyper-parameter which is easily interpretable. As such it is chosen as the best classifier for this problem. Since RF is an ensemble technique, out-of-bag (OOB) trees can be used to provide another estimate for the effectiveness measure [1].

Learning technique	Hyper-parameters	v_{effect}	s.d.
Decision tree	n_{min} : minimum number of samples required to be at a leaf node	0.63	0.062
Random forest	n_{tree} : number of trees in the forest	0.79	0.035
SVM Gaussian kernel	γ : kernel coefficient c : regularization parameter	0.79	0.059
SVM linear kernel	c : regularization parameter	0.80	0.059
KNN	k : number of neighbours d : type of distance	0.69	0.010

Table 1: Results of the nested cross-validation.

The final RF is fitted on the whole data and hyper-parameter tuning is performed using a grid-search approach. From this fitting procedure, the OOB estimates are obtained. After that the balanced accuracy and the TPR for each class are obtained using those estimates.

4.3 Results and discussion

With the final grid-search procedure, the selected number of trees is $n_{\text{tree}} = 500$. The balanced accuracy of the model is v_{effect} , obtained through the nested cross-validation, is 0.79. The same index is also computed using OOB estimation. In this case, a balanced accuracy of 0.76 is achieved. This is quite similar to the one obtained previously.

Through OOB error estimation, it is also possible to compute an estimate of which species our classifier correctly identifies and which it does not.

It appears that the classes which are misclassified the most are class 4 (i.e. *Alnus*) and class 32 (i.e. *Acca Sellowiana*). To ensure this was not dependent on the random component of the model, a set of different seeds are chosen and the results were always the same.

In terms of effectiveness, our classifier seems to perform quite well as shown by the balanced accuracy measure. However, the performance is poor on some classes as outlined above.

In terms of efficiency, RF are quite fast. For the prediction of the species of a new leaf (out of the ones present in the dataset) very little time is required, making this classifier a fast tool to help with leaf classification tasks.

Interpretability is not great. Variable importance could have been assessed to provide some local interpretability, but after noticing that the variables are correlated, it was not deemed appropriate.

A possible extension of this work would be to further examine the classes which are misclassified the most by our proposed classifier. To do this, oversampling could be tried out on the smallest classes, in order to detect if misclassification is due to some classes being less frequent or to inherent characteristics of those leaves. If the latter were the case, further features could be extracted from the images (this would require domain experts).

References

- [1] Leo Breiman. *Out-of-Bag estimation*. Technical report, Dept. Statistics, Univ. California, Berkeley, CA., 1997.
- [2] Margherita Grandini, Enrico Bagli, and Giorgio Visani. *Metrics for multi-class classification: an overview*. arXiv:2008.05756, 2020.
- [3] Akhilesh Gupta, Nesime Tatbul, Ryan Marcus, Shengtian Zhou, Insup Lee, and Justin Gottschlich. *Class-weighted evaluation metrics for imbalanced data classification*. arXiv:2010.05995, 2020.
- [4] Sebastian Raschka. *Model Evaluation, Model Selection, and Algorithm Selection in Machine Learning*. arXiv:1811.12808, 2018.
- [5] Pedro Filipe Silva. *Development of a System for Automatic Plant Species Recognition*. Master’s thesis, Faculdade de Ciencias da Universidade do Porto, 2013.
- [6] Fabian Spaeh and Sven Kosub. *Global Evaluation for Decision Tree Learning*. arXiv.2208.04828, 2022.